

ANALYTICAL SQL

DECEMBER 2021



Agenda

❖ Day 1

- Introduction to Analytical SQL
- Analytic window functions
- Ranking functions
- LAB

❖ Day 2

- Aggregate Analytical functions
- Windowing
- LAB

❖ Day 3

- Pivoting operations
- Statistical Aggregates
- Case Study

Introduction to Analytical SQL

- Example : With Vs. Without Analytical Function
- Why Analytical Functions?
- Aggregates vs. Analytics
- Syntax & Execution Sequence

Example : With Vs. Without Analytical Function

```
SELECT
    empno, ename, job, hiredate, sal
FROM emp;
```

EMPNO	ENAME	JOB	HIREDATE	SAL
7369	SMITH	CLERK	17-DEC-80	800
7499	ALLEN	SALESMAN	20-FEB-81	1600
7521	WARD	SALESMAN	22-FEB-81	1250
7566	JONES	MANAGER	02-APR-81	2975
7654	MARTIN	SALESMAN	28-SEP-81	1250
7698	BLAKE	MANAGER	01-MAY-81	2850
7782	CLARK	MANAGER	09-JUN-81	2450
7788	SCOTT	ANALYST	19-APR-87	3000
7839	KING	PRESIDENT	17-NOV-81	5000
7844	TURNER	SALESMAN	08-SEP-81	1500
7876	ADAMS	CLERK	23-MAY-87	1100

Example : With Vs. Without Analytical Function

The sequence in which everyone joined the company "HIRE_SEQ" Ordered by Salary

EMPNO	ENAME	JOB	HIREDATE	SAL	HIRE_SEQ
7369	SMITH	CLERK	17-DEC-80	800	1
7900	JAMES	CLERK	03-DEC-81	950	10
7876	ADAMS	CLERK	23-MAY-87	1100	14
7521	WARD	SALESMAN	22-FEB-81	1250	3
7654	MARTIN	SALESMAN	28-SEP-81	1250	8
7934	MILLER	CLERK	23-JAN-82	1300	12
7844	TURNER	SALESMAN	08-SEP-81	1500	7
7499	ALLEN	SALESMAN	20-FEB-81	1600	2
7782	CLARK	MANAGER	09-JUN-81	2450	6
7698	BLAKE	MANAGER	01-MAY-81	2850	5
7566	JONES	MANAGER	02-APR-81	2975	4

Example : With Vs. Without Analytical Function

Without Analytical Functions?

```
SELECT
    e.empno, e.ename, e.job, e.hiredate, e.sal, x.seq as hire_seq
FROM emp e ,
    (SELECT e2.empno, count(*) seq
     FROM emp e1, emp e2
     WHERE e1.hiredate <= e2.hiredate
     GROUP BY e2.empno
    ) x
WHERE e.empno = x.empno
ORDER BY sal;
```

OR

```
SELECT E2."EMPNO",E2."ENAME" , E2."JOB" , E2."HIREDATE"
    ,E2."SAL", COUNT(*) HIRE_SEQ

FROM "General_schema".emp E1
    , "General_schema".emp E2

WHERE E1."HIREDATE" <= E2."HIREDATE"
GROUP BY E2."EMPNO",E2."ENAME" , E2."JOB" , E2."HIREDATE",E2."SAL"
ORDER BY E2."SAL"
```

With Analytical Functions?

```
SELECT
    empno, ename, job, hiredate, sal,
    rank() over (order by hiredate) as hire_seq
FROM emp
ORDER BY sal;
```

Why Analytic Functions?

- Ability to see one row from another row's perspective
- Avoid self-join queries
- Summary data in detailed rows
- Slice and dice within the results

Aggregates vs. Analytics

- Aggregate functions
 - Rows are collapsed. One row per group
 - Group-BY columns must exist in the SELECT list
- Analytic functions
 - Rows are not collapsed
 - As many rows in the output as in the input
 - No restrictions on the columns in the SELECT list
 - Can appear only in the SELECT or ORDER BY clause
 - Evaluated after joins, WHERE, GROUP BY, HAVING clauses

Aggregates vs. Analytics (Example)

```
SELECT empno, deptno, sal,  
       AVG(sal) OVER () AS avg_sal  
FROM emp;
```

EMPNO	DEPTNO	SAL	AVG_SAL
7369	20	800	2073.21429
7499	30	1600	2073.21429
7521	30	1250	2073.21429
7566	20	2975	2073.21429
7654	30	1250	2073.21429
7698	30	2850	2073.21429
7782	10	2450	2073.21429
7788	20	3000	2073.21429
7839	10	5000	2073.21429
7844	30	1500	2073.21429
7876	20	1100	2073.21429

```
SELECT  
       AVG(sal) as avg_sal  
FROM emp;
```

```
       AVG_SAL  
-----  
2073.21429
```

Simplified Syntax

FUNCTION(<arg>,<arg>,...)

OVER (

<partition clause>

<sorting clause>

<windowing clause>

)

Execution Sequence

1. Table Joins
2. WHERE clause filters
3. GROUP BY
4. HAVING
5. Analytic Functions
6. DISTINCT
7. ORDER BY

Analytical Window Functions

- OVER Clause
- PARTITION BY Clause
- ORDER BY Clause
 - LEAD/LAG Functions

Over clause

The OVER clause is used to determine

when the function's calculations should restart ==> (PARTITION BY)

what order they are evaluated in by that function ==> (ORDER BY)

which rows from the query are applied to the function ==> (ROWS or RANGE)

```
<function> OVER (
    [PARTITION BY clause]
    [ORDER BY clause]
    [ROWS or RANGE clause])
```

In looking at the syntax, it appears that all of the sub-clauses are optional. In fact, each function that can use the OVER clause determines which of the sub-clauses are allowed, and which are required. Depending on the function being used, the OVER clause itself may be optional ([more details about the functions will be discussed later](#)).

Calculate the average salary per department

```
SELECT
    deptno, AVG(sal) as avg_sal
FROM emp
group by deptno;
```

DEPTNO	AVG_SAL
-----	-----
30	1566.66667
20	2175
10	2916.66667

One record for each
group

OVER (PARTITION BY..)

```
SELECT empno, deptno, sal,  
       AVG(sal) OVER (PARTITION BY deptno) AS avg_dept_sal  
FROM   emp;
```

One result for each
record in the dataset. No
grouping

EMPNO	DEPTNO	SAL	AVG_DEPT_SAL
7782	10	2450	2916.66667
7839	10	5000	2916.66667
7934	10	1300	2916.66667
7566	20	2975	2175
7902	20	3000	2175
7876	20	1100	2175
7369	20	800	2175
7788	20	3000	2175
7521	30	1250	1566.66667
7844	30	1500	1566.66667
7499	30	1600	1566.66667

PARTITION BY Clause

- The `query_partition_clause` divides the result set into partitions, or groups of data.
- The operation of the analytic function is restricted to the boundary imposed by these partitions, similar to the way a `GROUP BY` clause affects the action of an aggregate function.
- If the `query_partition_clause` is omitted, the whole result set is treated as a single partition.

HANDS ON

Calculate the number of employees per job

DEPTNO	ENAME	SAL	JOB	JOBCOUNT
20	SCOTT	3000	ANALYST	2
20	FORD	3000	ANALYST	2
10	MILLER	1300	CLERK	4
30	JAMES	950	CLERK	4
20	SMITH	800	CLERK	4
20	ADAMS	1100	CLERK	4
30	BLAKE	2850	MANAGER	3
20	JONES	2975	MANAGER	3
10	CLARK	2450	MANAGER	3
10	KING	5000	PRESIDENT	1
30	TURNER	1500	SALESMAN	4

HANDS ON

- With no analytical function:

```
SELECT deptno, ename, sal, job
, ( SELECT COUNT ( * ) FROM emp WHERE job = e.job ) jobcount
FROM emp e;
```

- With analytical function:

```
SELECT deptno, ename, sal, job
, COUNT ( * ) OVER ( PARTITION BY job ) jobcount
FROM emp;
```

ORDER BY Clause

- The query_order_clause controls the order that the rows are evaluated by the function
- The query_order_clause defines the logical order of the rows within each partition of the result set.
- ASC | DESC
Specifies that the values in the specified column should be sorted in ascending or descending order. ASC is the default sort order.

LEAD and LAG Functions

- Return the value from a field when looking one record (or more) behind/ahead using the order specified
- Syntax : LAG (field_name, <num_recs>, <default_value>) OVER (**ORDER BY** field_name)
- ORDER BY is **required**
- Optional second parameter to look more than one record

LEAD and LAG Functions (Example)

```
SELECT *,  
LEAD("SAL",1) OVER(ORDER BY "SAL") AS Next_Higher_Salary,  
LAG("SAL",1,0) OVER(ORDER BY "SAL") AS PREV_LOWER_Salary
```

```
From "General_schema".emp
```

EMPNO integer	ENAME character varying	JOB character varying	HIREDATE date	SAL numeric	next_higher_salary numeric	prev_lower_salary numeric
7369	SMITH	CLERK	1980-12-17	800	1100	0
7876	ADAMS	CLERK	1987-05-23	1100	1250	800
7521	WARD	SALESMAN	1981-02-22	1250	1250	1100
7654	MARTIN	SALESMAN	1981-09-28	1250	1500	1250
7844	TURNER	SALESMAN	1981-09-08	1500	1600	1250
7499	ALLEN	SALESMAN	1981-02-20	1600	2450	1500
7782	CLARK	MANAGER	1981-06-09	2450	2850	1600
7698	BLAKE	MANAGER	1981-05-01	2850	2975	2450
7566	JONES	MANAGER	1981-04-02	2975	3000	2850
7788	SCOTT	ANALYST	1987-04-19	3000	5000	2975
7839	KING	PRESIDENT	1981-11-17	5000	[null]	3000

HANDS ON

Get the previous hire date for each employee then calculate the hiring gap in days

EMPNO	ENAME	HIREDATE	PREV_HIRE_DATE	Hiring_Gap
7369	SMITH	1980-12-17	NULL	NULL
7499	ALLEN	1981-02-20	1980-12-17	65
7521	WARD	1981-02-22	1981-02-20	2

HANDS ON

```
SELECT *, "HIREDATE" - X."PREV_HIRE_DATE" AS Hiring_Gap
FROM ( select "EMPNO" , "ENAME" , "HIREDATE",
      LAG("HIREDATE",1) OVER(ORDER BY "HIREDATE") AS "PREV_HIRE_DATE"
from "General_schema".emp ) AS x
```

OR

```
SELECT "EMPNO" , "ENAME" , "HIREDATE",
      LAG("HIREDATE",1) OVER(ORDER BY "HIREDATE") AS "PREV_HIRE_DATE",
      "HIREDATE" - LAG("HIREDATE",1) OVER(ORDER BY "HIREDATE") AS Hiring_Gap
FROM "General_schema".emp
```

PARTITION BY & ORDER BY

```
SELECT deptno, empno, sal
, LEAD ( sal, 1, 0 ) OVER
  ( PARTITION BY deptno ORDER BY sal DESC NULLS LAST ) next_lower_sal
, LAG ( sal, 1, 0 ) OVER
  ( PARTITION BY deptno ORDER BY sal DESC NULLS LAST ) prev_higher_sal
FROM emp
WHERE deptno in ( 10, 20 )
ORDER BY deptno, sal DESC;
```

DEPTNO	EMPNO	SAL	NEXT_LOWER_SAL	PREV_HIGHER_SAL
10	7839	5000	2450	0
10	7782	2450	1300	5000
10	7934	1300	0	2450
20	7788	3000	3000	0
20	7902	3000	2975	3000
20	7566	2975	1100	3000
20	7876	1100	800	2975
20	7369	800	0	1100

Order Of Items In Analytic Clause

```
SELECT deptno, empno, ename, sal
, MIN ( sal ) OVER ( ORDER BY ename PARTITION BY deptno ) minsal
FROM emp;
```

```
Error at Command Line : 2 Column : 37
Error report -
SQL Error: ORA-00907: missing right parenthesis
00907. 00000 - "missing right parenthesis"
```

Components must be in
correct order

Ranking Functions

- RANK , DENSE_RANK and ROW_NUMBER Function
- FIRST_VALUE and LAST_VALUE Function
- PERCENT_RANK Function
- NTILE Function
- CUME_DIST Function

Ranking Functions

- Where does this record fall, when the records are placed in a certain order?
- Ordering (Ranking) functions:
 - RANK
 - DENSE_RANK
 - ROW_NUMBER
- Syntax:
 - RANK () OVER (ORDER BY field_name)
- ORDER BY expression is **mandatory** for Ranking function
- All three functions return a number (Rank)
- Difference between functions is how they handle ties

```

SELECT deptno, ename, sal
, RANK () OVER ( ORDER BY ename ) r1
, DENSE_RANK () OVER ( ORDER BY ename ) r2
, ROW_NUMBER () OVER ( ORDER BY ename ) r3
FROM emp
ORDER BY ename;

```

DEPTNO	ENAME	SAL	R1	R2	R3
20	ADAMS	1100	1	1	1
30	ALLEN	1600	2	2	2
30	BLAKE	2850	3	3	3
10	CLARK	2450	4	4	4
20	FORD	3000	5	5	5
30	JAMES	950	6	6	6
20	JONES	2975	7	7	7
10	KING	5000	8	8	8
30	MARTIN	1250	9	9	9
10	MILLER	1300	10	10	10
20	SCOTT	3000	11	11	11

When there are no ties,
all three of these
functions return the
same values.

```

SELECT ename, sal
, RANK ( ) OVER ( ORDER BY sal ) r1
, DENSE_RANK ( ) OVER ( ORDER BY sal ) r2
, ROW_NUMBER ( ) OVER ( ORDER BY sal ) r3
FROM emp
ORDER BY sal;

```

ENAME	SAL	R1	R2	R3
SMITH	800	1	1	1
JAMES	950	2	2	2
ADAMS	1100	3	3	3
WARD	1250	4	4	4
MARTIN	1250	4	4	5
MILLER	1300	6	5	6
TURNER	1500	7	6	7
ALLEN	1600	8	7	8
CLARK	2450	9	8	9
BLAKE	2850	10	9	10
JONES	2975	11	10	11

RANK and DENSE_RANK will assign the same number to multiple records with the same sort value

The difference is in how each one handles the record which follows

ROW_NUMBER assigns a unique number to each record. The highest value assigned by ROW_NUMBER will be equal to COUNT(*)

Find out the highest salary per department

- Step 1:

```
SELECT empno, deptno, sal,  
       RANK() OVER (PARTITION BY deptno ORDER BY sal DESC) AS rnk  
FROM emp;
```

EMPNO	DEPTNO	SAL	RNK
7839	10	5000	1
7782	10	2450	2
7934	10	1300	3
7788	20	3000	1
7902	20	3000	1
7566	20	2975	3
7876	20	1100	4
7369	20	800	5
7698	30	2850	1
7499	30	1600	2
7844	30	1500	3

Find out the highest salary per department

- Step 2:

```
SELECT
  rnk_sal.empno as Empolyee_id
, rnk_sal.deptno as Department_id
, rnk_sal.sal as Highest_salary
FROM
  (
    SELECT empno, deptno, sal,
           RANK() OVER
             (PARTITION BY deptno ORDER BY sal DESC) AS rnk
    FROM emp
  ) rnk_sal
WHERE rnk_sal.rnk = 1;
```

EMPLOYEE_ID	DEPARTMENT_ID	HIGHEST_SALARY
-----	-----	-----
7839	10	5000
7788	20	3000
7902	20	3000
7698	30	2850

Find out the highest salary per department

- Step 1:

```
SELECT empno, deptno, sal,  
       DENSE_RANK() OVER (PARTITION BY deptno ORDER BY sal DESC) AS rnk  
FROM emp;
```

EMPNO	DEPTNO	SAL	RNK
7839	10	5000	1
7782	10	2450	2
7934	10	1300	3
7788	20	3000	1
7902	20	3000	1
7566	20	2975	2
7876	20	1100	3
7369	20	800	4
7698	30	2850	1
7499	30	1600	2
7844	30	1500	3

Find out the highest salary per department

- Step 2:

```
SELECT
  rnk_sal.empno as Empolyee_id
, rnk_sal.deptno as Department_id
, rnk_sal.sal as Highest_salary
FROM
  (
    SELECT empno, deptno, sal,
           DENSE_RANK() OVER
             (PARTITION BY deptno ORDER BY sal DESC) AS rnk
    FROM emp
  ) rnk_sal
WHERE rnk_sal.rnk = 1;
```

EMPLOYEE_ID	DEPARTMENT_ID	HIGHEST_SALARY
-----	-----	-----
7839	10	5000
7788	20	3000
7902	20	3000
7698	30	2850

Find out the highest salary per department

- Step 1:

```
SELECT empno, deptno, sal,  
       ROW_NUMBER() OVER (PARTITION BY deptno ORDER BY sal DESC) AS rnk  
FROM emp;
```

EMPNO	DEPTNO	SAL	RNK
7839	10	5000	1
7782	10	2450	2
7934	10	1300	3
7788	20	3000	1
7902	20	3000	2
7566	20	2975	3
7876	20	1100	4
7369	20	800	5
7698	30	2850	1
7499	30	1600	2
7844	30	1500	3

Find out the highest salary per department

- Step 2:


```
SELECT
  rnk_sal.empno as Empolyee_id
,rnk_sal.deptno as Department_id
,rnk_sal.sal as Highest_salary
FROM
  (
    SELECT empno, deptno, sal,
           ROW_NUMBER() OVER
             (PARTITION BY deptno ORDER BY sal DESC) AS rnk
    FROM emp
  ) rnk_sal
WHERE rnk_sal.rnk = 1;
```

EMPOLYEE_ID	DEPARTMENT_ID	HIGHEST_SALARY
-----	-----	-----
7839	10	5000
7788	20	3000
7698	30	2850

HANDS ON

rank each host based on the number of beds they have listed on our website. The host with the most beds should be ranked first (rank = 1)

host_id	apartment_id	apartment_type	n_beds	n_bedrooms	city
0	A1	Room	1	1	New York
0	A2	Room	1	1	New Jersey
0	A3	Room	1	1	New Jersey
1	A4	Apartment	2	1	Houston
1	A5	Apartment	2	1	Las Vegas
2	A6	Yurt	3	1	-
3	A7	Penthouse	3	3	Tianjin
3	A8	Penthouse	5	5	Beijing



host_id	number_of_beds	rank
10	16	1
3	8	2
6	6	3
5	5	4
7	4	5
1	4	5
9	4	5
0	3	6
2	3	6
8	2	7
4	2	7
11	2	7

HANDS ON

rank each host based on the number of beds they have listed on our website. The host with the most beds should be ranked first (rank = 1)

```
SELECT
host_id,
sum(n_beds) as number_of_beds,
DENSE_RANK() OVER(ORDER BY sum(n_beds) DESC) as rank
FROM airbnb_apartments
GROUP BY 1
```

FIRST_VALUE and LAST_VALUE

- FIRST_VALUE / LAST_VALUE { (expr [{RESPECT | IGNORE} NULLS]) }
OVER (analytic_clause)
- Allows you to return the first / last result from an ordered set.
- The "{RESPECT | IGNORE} NULLS" clause indicates if NULLs are considered when determining results.

```

SELECT empno, deptno, sal,
       FIRST_VALUE( sal IGNORE NULLS )
         OVER (PARTITION BY deptno ORDER BY sal) AS lowest_in_dept
FROM   emp
ORDER BY deptno, sal;

```

EMPNO	DEPTNO	SAL	LOWEST_IN_DEPT
-----	-----	-----	-----
7934	10	1300	1300
7782	10	2450	1300
7839	10	5000	1300
7369	20	800	800
7876	20	1100	800
7566	20	2975	800
7788	20	3000	800
7902	20	3000	800
7900	30	950	950
7654	30	1250	950
7521	30	1250	950

FIRST_VALUE and LAST_VALUE (IGNORE NULLS)

- Though **LAST_VALUE** and **FIRST_VALUE** are quite handy on many occasions, one of the main limitations for the functions have been the absence of **IGNORE_NULLS** support, as found in many other RDBMS.

```
SELECT "ENAME", "EMPNO" , "JOB" ,"SAL"  
,FIRST_VALUE("SAL") OVER(PARTITION BY "JOB" ORDER BY "SAL" desc) AS "HIGHEST_IN_TEAM"  
FROM "General_schema".emp  
ORDER BY "JOB" , "SAL"
```

EMPNO	JOB	SAL	HIGHEST_IN_TEAM
7369	CLERK	800	[null]
7876	CLERK	1100	[null]
78732	CLERK	[null]	[null]
7521	SALESMAN	1250	[null]
7654	SALESMAN	1250	[null]
7844	SALESMAN	1500	[null]
7499	SALESMAN	1600	[null]
78739	SALESMAN	[null]	[null]

FIRST_VALUE and LAST_VALUE (IGNORE NULLS)

```
SELECT "EMPNO" , "JOB" ,"SAL"  
,FIRST_VALUE("SAL") OVER(PARTITION BY "JOB" ORDER BY  
    case when "SAL" is not null then "SAL"  
    else 0 end desc) AS "HIGHEST_IN_TEAM"  
  
FROM "General_schema".emp  
WHERE "JOB" IN ('CLERK' , 'SALESMAN')  
ORDER BY "JOB" , "SAL"
```

EMPNO	JOB	SAL	HIGHEST_IN_DEPT
7369	CLERK	800	1100
7876	CLERK	1100	1100
78732	CLERK	[null]	1100
7654	SALESMAN	1250	1600
7521	SALESMAN	1250	1600
7844	SALESMAN	1500	1600
7499	SALESMAN	1600	1600
78739	SALESMAN	[null]	1600

Date of next order(LAB)

- From orders table in the order entry schema (OE), Find the date of next order for each customer.
- There are almost 5 different ways to calculate the date of the next order, get as much as you can!

```
SELECT customer_id
, TRUNC ( order_date ) AS order_date
, order_total
, LEAD ( TRUNC(order_date) ) OVER
    ( PARTITION BY customer_id ORDER BY order_date ) AS next_order_date_LEAD
, LAG ( TRUNC(order_date) ) OVER
    ( PARTITION BY customer_id ORDER BY order_date DESC ) AS next_order_date_LAG
, MAX ( TRUNC(order_date) ) OVER
    ( PARTITION BY customer_id ORDER BY order_date
      ROWS BETWEEN CURRENT ROW AND 1 FOLLOWING ) AS next_order_date_MAX
, MIN ( TRUNC(order_date) ) OVER
    ( PARTITION BY customer_id ORDER BY order_date
      ROWS BETWEEN 1 FOLLOWING AND 1 FOLLOWING ) AS next_order_date_MIN
, MIN ( TRUNC(order_date) ) OVER
    ( PARTITION BY customer_id ORDER BY order_date
      ROWS BETWEEN 1 FOLLOWING AND UNBOUNDED FOLLOWING ) AS next_order_date_MIN2
FROM orders
ORDER BY 1, 2;
```